

# Boosting

Victor Kitov  
v.v.kitov@yandex.ru

Yandex School of Data Analysis



# Table of Contents

- 1 Forward stagewise additive modeling
- 2 Gradient boosting

# Linear ensembles

## Linear ensemble:

$$F(x) = f_0(x) + c_1 h_1(x) + \dots + c_M h_M(x)$$

**Regression:**  $\hat{y}(x) = F(x)$

**Binary classification:**  $\text{score}(y|x) = F(x)$ ,  $\hat{y}(x) = \text{sign } F(x)$

- Notation:  $h_1(x), \dots, h_M(x)$  are called *base learners*, *weak learners*, *base models*.
- Too expensive to optimize  $f_0(x), h_1(x), \dots, h_M(x)$  and  $c_1, \dots, c_M$  jointly for large  $M$ .
- Idea: optimize  $f_0(x)$  and then each pair  $(h_m(x), c_m)$  greedily.
- After ensemble is built we can fine-tune  $c_1, \dots, c_M$  by fitting features  $f_0(x), h_1(x), \dots, h_M(x)$  with linear regression/classifier.

# Forward stagewise additive modeling (FSAM)

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$ , general form of “base learner”  $h(x|\gamma)$  (dependent from parameter  $\gamma$ ) and the number  $M$  of successive additive approximations.

① Fit initial approximation  $f_0(x) = \arg \min_f \sum_{i=1}^N L(f(x_i), y_i)$

② For  $m = 1, 2, \dots, M$ :

① find next best classifier

$$(c_m, h_m) = \arg \min_{h, c} \sum_{i=1}^N L(f_{m-1}(x_i) + ch(x_i), y_i)$$

② set

$$f_m(x) = f_{m-1}(x) + c_m h_m(x)$$

**Output:** approximation function

$$f_M(x) = f_0(x) + \sum_{m=1}^M c_m h_m(x)$$

# Comments on FSAM

- Number of steps  $M$  should be determined by performance on validation set.
- Step 1 need not be solved accurately, since its mistakes are expected to be corrected by future base learners.
  - we can take  $f_0(x) = \arg \min_{\beta \in \mathbb{R}} \sum_{i=1}^N L(\beta, y_i)$  or simply  $f_0(x) \equiv 0$ .
- By similar reasoning there is no need to solve 2.1 accurately
  - typically very simple base learners are used such as trees of depth=1,2,3.
- For some loss functions, such as  $L(y, f(x)) = e^{-yf(x)}$  we can solve FSAM explicitly.
- For general loss functions gradient boosting scheme should be used.

## Adaboost (discrete version): assumptions

- binary classification task  $y \in \{+1, -1\}$
- family of base classifiers  $h(x) = h(x|\gamma)$  where  $\gamma$  is some fitted parametrization.
- $h(x) \in \{+1, -1\}$
- classification is performed with
$$\hat{y} = \text{sign}\{f_0(x) + c_1 f_1(x) + \dots + c_M f_M(x)\}$$
- optimized loss is  $L(y, f(x)) = e^{-yf(x)}$
- FSAM is applied

# Adaboost (discrete version): algorithm

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; number of additive weak classifiers  $M$ , a family of weak classifiers  $h(x) \in \{+1, -1\}$ , trainable on weighted datasets.

- 1 Initialize observation weights  $w_i = 1/n$ ,  $i = 1, 2, \dots, n$ .
- 2 for  $m = 1, 2, \dots, M$ :
  - 1 fit  $h^m(x)$  to training data using weights  $w_i$
  - 2 compute weighted misclassification rate:

$$E_m = \frac{\sum_{i=1}^N w_i \mathbb{I}[h^m(x) \neq y_i]}{\sum_{i=1}^N w_i}$$

- 3 if  $E_m > 0.5$  or  $E_m = 0$ : terminate procedure.
- 4 compute  $\alpha_m = \ln((1 - E_m)/E_m)$
- 5 increase all weights, where misclassification with  $h^m(x)$  was made:

$$w_i \leftarrow w_i e^{\alpha_m}, i \in \{i : h^m(x_i) \neq y_i\}$$

**Output:** composite classifier  $f(x) = \text{sign} \left( \sum_{m=1}^M \alpha_m h^m(x) \right)$

# Table of Contents

- 1 Forward stagewise additive modeling
- 2 Gradient boosting



# Motivation

- Problem: For general loss function  $L$  FSAM cannot be solved explicitly
- Analogy with function minimization: when we can't find optimum explicitly we use numerical methods
- Gradient boosting: numerical method for iterative loss minimization

# Gradient descent algorithm

$$F(w) \rightarrow \min_w, \quad w \in \mathbb{R}^N$$

Gradient descend algorithm:

**INPUT:**

$\eta$ -parameter, controlling the speed of convergence

$M$ -number of iterations

**ALGORITHM:**

initialize  $w$

**for**  $m = 1, 2, \dots, M$ :

$$\Delta w = \frac{\partial F(w)}{\partial w}$$

$$w = w - \eta \Delta w$$

# Modified gradient descent algorithm

**INPUT:**

$M$ -number of iterations

**ALGORITHM:**

initialize  $w$

**for**  $m = 1, 2, \dots M$ :

$$\Delta w = \frac{\partial F(w)}{\partial w}$$

$$c^* = \arg \min_c F(w - c\Delta w)$$

$$w = w - c^* \Delta w$$

# Gradient boosting

- Now consider  $F(f(x_1), \dots, f(x_N)) = \sum_{n=1}^N L(f(x_n), y_n)$
- Gradient descent performs pointwise optimization, but we need generalization, so we optimize in space of functions.
- Gradient boosting implements modified gradient descent in function space:
  - find  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - fit base learner  $h_m(x)$  to  $\{(x_i, z_i)\}_{i=1}^N$

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- 2 For each step  $m = 1, 2, \dots, M$ :

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- 2 For each step  $m = 1, 2, \dots, M$ :
  - 1 calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- 2 For each step  $m = 1, 2, \dots, M$ :
  - 1 calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$
  - 2 fit  $h_m$  to  $\{(x_i, z_i)\}_{i=1}^N$ , for example by solving

$$\sum_{n=1}^N (h_m(x_n) - z_n)^2 \rightarrow \min_{h_m}$$



# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$
  - ② fit  $h_m$  to  $\{(x_i, z_i)\}_{i=1}^N$ , for example by solving

$$\sum_{n=1}^N (h_m(x_n) - z_n)^2 \rightarrow \min_{h_m}$$

- ③ solve univariate optimization problem:

$$\sum_{i=1}^N L(f_{m-1}(x_i) + c_m h_m(x_i), y_i) \rightarrow \min_{c_m \in \mathbb{R}_+}$$

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$
  - ② fit  $h_m$  to  $\{(x_i, z_i)\}_{i=1}^N$ , for example by solving

$$\sum_{n=1}^N (h_m(x_n) - z_n)^2 \rightarrow \min_{h_m}$$

- ③ solve univariate optimization problem:

$$\sum_{i=1}^N L(f_{m-1}(x_i) + c_m h_m(x_i), y_i) \rightarrow \min_{c_m \in \mathbb{R}_+}$$

- ④ set  $f_m(x) = f_{m-1}(x) + c_m h_m(x)$

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- 2 For each step  $m = 1, 2, \dots, M$ :
  - 1 calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$
  - 2 fit  $h_m$  to  $\{(x_i, z_i)\}_{i=1}^N$ , for example by solving

$$\sum_{n=1}^N (h_m(x_n) - z_n)^2 \rightarrow \min_{h_m}$$

- 3 solve univariate optimization problem:

$$\sum_{i=1}^N L(f_{m-1}(x_i) + c_m h_m(x_i), y_i) \rightarrow \min_{c_m \in \mathbb{R}_+}$$

- 4 set  $f_m(x) = f_{m-1}(x) + c_m h_m(x)$

**Output:** approximation function  $f_M(x) = f_0(x) + \sum_{m=1}^M c_m h_m(x)$

# Gradient boosting: examples

In gradient boosting

$$\sum_{n=1}^N \left( h_m(x_n) - \left( -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x_n)} \right) \right)^2 \rightarrow \min_{h_m}$$

Specific cases:

- $L = \frac{1}{2} (r - y)^2 \Rightarrow -\frac{\partial L}{\partial r} = -(r - y) = (y - r)$ 
  - $h_m(x)$  is fitted to compensate regression errors  $(y - f_{m-1}(x))$
- $L = [-ry]_+$
- $L = \ln(1 + e^{-ry})$

## Modification of boosting for trees

- Compared to first method of gradient boosting, boosting of regression trees finds additive coefficients individually for each terminal region  $R_{jm}$ , not globally for the whole classifier  $h^m(x)$ .
- This is done to increase accuracy: forward stagewise algorithm cannot be applied to find  $R_{jm}$ , but it can be applied to find  $\gamma_{jm}$ , because second task is solvable for arbitrary  $L$ .
- Max leaves  $J$ 
  - interaction between no more than  $J - 1$  terms
  - usually  $J \leq 8$

## Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit constant initial approximation  $f_0(x)$ :

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(\gamma, y_i)$$

## Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit constant initial approximation  $f_0(x)$ :

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(\gamma, y_i)$$

- 2 For each step  $m = 1, 2, \dots, M$ :

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit constant initial approximation  $f_0(x)$ :

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(\gamma, y_i)$$

- 2 For each step  $m = 1, 2, \dots, M$ :

- 1 calculate derivatives  $\mathbf{z}_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$



# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit constant initial approximation  $f_0(x)$ :

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(\gamma, y_i)$$

- 2 For each step  $m = 1, 2, \dots, M$ :

- 1 calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
- 2 fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^N$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit constant initial approximation  $f_0(x)$ :  

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(\gamma, y_i)$$
- 2 For each step  $m = 1, 2, \dots, M$ :
  - 1 calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - 2 fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^N$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .
  - 3 for each terminal region  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(f_{m-1}(x_i) + \gamma, y_i)$$

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit constant initial approximation  $f_0(x)$ :  

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(\gamma, y_i)$$
- 2 For each step  $m = 1, 2, \dots, M$ :
  - 1 calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f_{m-1}(x)}$
  - 2 fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^N$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .
  - 3 for each terminal region  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(f_{m-1}(x_i) + \gamma, y_i)$$

- 4 update  $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} \mathbb{I}[x \in R_{jm}]$

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit constant initial approximation  $f_0(x)$ :  

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(\gamma, y_i)$$
- 2 For each step  $m = 1, 2, \dots, M$ :
  - 1 calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - 2 fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^N$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .
  - 3 for each terminal region  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(f_{m-1}(x_i) + \gamma, y_i)$$

- 4 update  $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} \mathbb{I}[x \in R_{jm}]$

**Output:** approximation function  $f_M(x)$

## Shrinkage & subsampling

- Shrinkage of general GB, step (d):

$$f_m(x) = f_{m-1}(x) + \nu c_m h_m(x)$$

- Shrinkage of trees GB, step (d):

$$f_m(x) = f_{m-1}(x) + \nu \sum_{j=1}^{J_m} \gamma_{jm} \mathbb{I}[x \in R_{jm}]$$

- Comments:

- $\nu \in (0, 1]$
- $\nu \downarrow \implies M \uparrow$

- Subsampling

- increases speed of fitting
- may increase accuracy

## Case of $C \geq 3$ classes

- Can fit  $C$  independent boostings (one vs. all scheme)
  - $\hat{y} = \arg \max_y f_{my}(x)$
- Alternatively can optimize multivariate  $L(f(x), y) = -\ln p(y|x)$ 
  - using linear or quadratic approximation
  - for quadratic approximation need to invert  $\frac{\partial^2 F(r, y)}{\partial r^2} \Big|_{r=f(x)}$ .  
Can use diagonal approximation.

# Types of boosting

- Loss function  $F$ :
  - $F(|f(x) - y|)$  - regression
  - $-\ln p(y|x)$  or  $F(y \cdot \text{score}(y = +1|x))$  - binary classification
- Optimization
  - analytical (AdaBoost)
  - gradient based
- Base learners
  - continuous
  - discrete
- Classification
  - binary
  - multiclass
- Extensions: shrinkage, subsampling